

# ML Problem Set 1

2025

## 1 Problem Statement

Suppose you are the CEO of a restaurant franchise and are considering different cities for opening a new outlet. - You would like to expand your business to cities that may give your restaurant higher profits. - The chain already has restaurants in various cities and you have data for profits and populations from the cities. - You also have data on cities that are candidates for a new restaurant. - For these cities, you have the city population.

Can you use the data to help you identify which cities may potentially give your business higher profits?

## 2 Packages

Import all the packages that you will need during this assignment.

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import math
import copy
```

## 3 Dataset

You will start by loading the dataset for this task. The dataset contains two variables, `x_train` and `y_train` - `x_train` is the population of a city - `y_train` is the profit of a restaurant in that city. A negative value for profit indicates a loss.

- Both `X_train` and `y_train` are numpy arrays.

```
[ ]: # Load the data from the Excel file
df_x_train = pd.read_excel('PS1_data.xlsx', sheet_name='X_train')
df_y_train = pd.read_excel('PS1_data.xlsx', sheet_name='Y_train')

# Convert the DataFrames back to numpy arrays
x_train = df_x_train.values.flatten()
y_train = df_y_train.values.flatten()

print("Data has been loaded from Excel successfully.")
```

**Inspect the data** Before starting any task, it is useful to get more familiar with your dataset. - A good place to start is to just print out each variable and see what it contains.

The code below prints the variable `x_train` and the type of the variable.

```
[ ]: # print x_train
print("Type of x_train:",type(x_train))
print("First five elements of x_train are:\n", x_train[:5])
```

`x_train` is a numpy array that contains decimal values that are all greater than zero. - These values represent the city population times 10,000 - For example, 6.1101 means that the population for that city is 61,101

Now, let's print `y_train`

```
[ ]: # print y_train
print("Type of y_train:",type(y_train))
print("First five elements of y_train are:\n", y_train[:5])
```

Similarly, `y_train` is a numpy array that has decimal values, some negative, some positive. - These represent your restaurant's average monthly profits in each city, in units of \$10,000. - For example, 17.592 represents \$175,920 in average monthly profits for that city. - -2.6807 represents -\$26,807 in average monthly loss for that city.

**Check the dimensions of your variables** Another useful way to get familiar with your data is to view its dimensions.

Please print the shape of `x_train` and `y_train` and see how many training examples you have in your dataset.

```
[ ]: print ('The shape of x_train is:', x_train.shape)
print ('The shape of y_train is: ', y_train.shape)
print ('Number of training examples (m):', len(x_train))
```

The city population array has 97 data points, and the monthly average profits also has 97 data points. These are NumPy 1D arrays.

### 3.1 Visualize your data

It is often useful to understand the data by visualizing it. - For this dataset, you can use a scatter plot to visualize the data, since it has only two properties to plot (profit and population). - Many other problems that you will encounter in real life have more than two properties (for example, population, average household income, monthly profits, monthly sales). When you have more than two properties, you can still use a scatter plot to see the relationship between each pair of properties.

```
[ ]: # Create a scatter plot of the data. To change the markers to red "x",
# we used the 'marker' and 'c' parameters
plt.scatter(x_train, y_train, marker='x', c='r')

# Set the title
plt.title("Profits vs. Population per city")
```

```
# Set the y-axis label
plt.ylabel('Profit in $10,000')
# Set the x-axis label
plt.xlabel('Population of City in 10,000s')
plt.show()
```

Your goal is to build a linear regression model to fit this data. - With this model, you can then input a new city's population, and have the model estimate your restaurant's potential monthly profits for that city.

## 4 Exercise 1. Compute Cost

Gradient descent involves repeated steps to adjust the value of your parameter  $(w, b)$  to gradually get a smaller and smaller cost  $J(w, b)$ . - At each step of gradient descent, it will be helpful for you to monitor your progress by computing the cost  $J(w, b)$  as  $(w, b)$  gets updated. - In this section, you will implement a function to calculate  $J(w, b)$  so that you can check the progress of your gradient descent implementation.

**Cost function** As you may recall from the lecture, for one variable, the cost function for linear regression  $J(w, b)$  is defined as

$$J(w, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)})^2$$

- You can think of  $f_{w,b}(x^{(i)})$  as the model's prediction of your restaurant's profit, as opposed to  $y^{(i)}$ , which is the actual profit that is recorded in the data.
- $m$  is the number of training examples in the dataset

### Model prediction

- For linear regression with one variable, the prediction of the model  $f_{w,b}$  for an example  $x^{(i)}$  is represented as:

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b$$

This is the equation for a line, with an intercept  $b$  and a slope  $w$

**Implementation** Please complete the `compute_cost()` function below to compute the cost  $J(w, b)$ .

### 4.1 Exercise

Complete the `compute_cost` below to:

- Iterate over the training examples, and for each example, compute:
  - The prediction of the model for that example

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b$$

- The cost for that example

$$cost^{(i)} = (f_{wb} - y^{(i)})^2$$

- Return the total cost over all examples

$$J(\mathbf{w}, b) = \frac{1}{2m} \sum_{i=0}^{m-1} cost^{(i)}$$

- Here,  $m$  is the number of training examples and  $\sum$  is the summation operator

```
[ ]: def compute_cost(x, y, w, b):
    """
    Computes the cost function for linear regression.

    Args:
        x (ndarray): Shape (m,) Input to the model (Population of cities)
        y (ndarray): Shape (m,) Label (Actual profits for the cities)
        w, b (scalar): Parameters of the model

    Returns
        total_cost (float): The cost of using w,b as the parameters for linear_
        ↪ regression
                               to fit the data points in x and y
    """
    # number of training examples
    m = x.shape[0] ## uncomment this line of code

    # You need to return this variable correctly
    total_cost = 0 ## uncomment this line of code

    ### START YOUR CODE HERE ###

    ### END YOUR CODE HERE ###

    return total_cost ## uncomment this line of code
```

```
[ ]: # Compute cost with some initial values for parameters w, b
initial_w = 2
initial_b = 1

cost = compute_cost(x_train, y_train, initial_w, initial_b)

print(f'Cost at initial w: {cost:.3f}')
```

**Expected Output:**

Cost at initial w: 75.203

## 5 Exercise 2. Gradient descent

In this section, you will implement the gradient for parameters  $w, b$  for linear regression.

As described in the lecture videos, the gradient descent algorithm is:

*Repeat until convergence:*

$$b := b - \alpha \frac{\partial J(w, b)}{\partial b} \quad w := w - \alpha \frac{\partial J(w, b)}{\partial w} \quad (1)$$

where, parameters  $w, b$  are both updated simultaneously and where

$$\frac{\partial J(w, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)}) \quad (2)$$

$$\frac{\partial J(w, b)}{\partial w} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)} \quad (3)$$

\*  $m$  is the number of training examples in the dataset

- $f_{w,b}(x^{(i)})$  is the model's prediction, while  $y^{(i)}$ , is the target value

You will implement a function called `compute_gradient` which calculates  $\frac{\partial J(w)}{\partial w}$ ,  $\frac{\partial J(w)}{\partial b}$

### 5.1 Exercise 2

Please complete the `compute_gradient` function to:

- Iterate over the training examples, and for each example, compute:
  - The prediction of the model for that example

$$f_{wb}(x^{(i)}) = wx^{(i)} + b$$

- The gradient for the parameters  $w, b$  from that example

$$\frac{\partial J(w, b)}{\partial b}^{(i)} = (f_{w,b}(x^{(i)}) - y^{(i)})$$

$$\frac{\partial J(w, b)}{\partial w}^{(i)} = (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$$

- Return the total gradient update from all the examples

$$\frac{\partial J(w, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} \frac{\partial J(w, b)}{\partial b}^{(i)}$$

$$\frac{\partial J(w, b)}{\partial w} = \frac{1}{m} \sum_{i=0}^{m-1} \frac{\partial J(w, b)}{\partial w}^{(i)}$$

- Here,  $m$  is the number of training examples and  $\sum$  is the summation operator

If you get stuck, you can check out the hints presented after the cell below to help you with the implementation.

```
[ ]: def compute_gradient(x, y, w, b):
    """
    Computes the gradient for linear regression
    Args:
        x (ndarray): Shape (m,) Input to the model (Population of cities)
        y (ndarray): Shape (m,) Label (Actual profits for the cities)
        w, b (scalar): Parameters of the model
    Returns
        dj_dw (scalar): The gradient of the cost w.r.t. the parameters w
        dj_db (scalar): The gradient of the cost w.r.t. the parameter b
    """

    # Number of training examples
    m = x.shape[0]

    # You need to return the following variables correctly
    dj_dw = 0
    dj_db = 0

    ### START YOUR CODE HERE ###

    ### END YOUR CODE HERE ###

    return dj_dw, dj_db

[ ]: # Compute and display cost and gradient with non-zero w
test_w = 0.2
test_b = 0.2
tmp_dj_dw, tmp_dj_db = compute_gradient(x_train, y_train, test_w, test_b)

print('Gradient at test w, b:', tmp_dj_dw, tmp_dj_db)
```

**Expected Output:**

Gradient at test w

-47.41610118 -4.007175051546391

## 6 Learning parameters using batch gradient descent

You will now find the optimal parameters of a linear regression model by using batch gradient descent. Recall batch refers to running all the examples in one iteration. - You don't need to implement anything for this part. Simply run the cells below.

- A good way to verify that gradient descent is working correctly is to look at the value of  $J(w, b)$  and check that it is decreasing with each step.
- Assuming you have implemented the gradient and computed the cost correctly and you have

an appropriate value for the learning rate  $\alpha$ ,  $J(w, b)$  should never increase and should converge to a steady value by the end of the algorithm.

```
[ ]: def gradient_descent(x, y, w_in, b_in, cost_function, gradient_function, alpha, num_iters):  
    """  
    Performs batch gradient descent to learn theta. Updates theta by taking  
    num_iters gradient steps with learning rate alpha  
  
    Args:  
        x : (ndarray): Shape (m,)   
        y : (ndarray): Shape (m,)   
        w_in, b_in : (scalar) Initial values of parameters of the model  
        cost_function: function to compute cost  
        gradient_function: function to compute the gradient  
        alpha : (float) Learning rate  
        num_iters : (int) number of iterations to run gradient descent  
    Returns  
        w : (ndarray): Shape (1,) Updated values of parameters of the model after  
            running gradient descent  
        b : (scalar) Updated value of parameter of the model after  
            running gradient descent  
    """  
  
    # number of training examples  
    m = len(x)  
  
    # An array to store cost J and w's at each iteration - primarily for  
    graphing later  
    J_history = []  
    w_history = []  
    w = copy.deepcopy(w_in) #avoid modifying global w within function  
    b = b_in  
  
    for i in range(num_iters):  
  
        # Calculate the gradient and update the parameters  
        dj_dw, dj_db = gradient_function(x, y, w, b )  
  
        # Update Parameters using w, b, alpha and gradient  
        w = w - alpha * dj_dw  
        b = b - alpha * dj_db  
  
        # Save cost J at each iteration  
        if i<100000: # prevent resource exhaustion  
            cost = cost_function(x, y, w, b)  
            J_history.append(cost)
```

```

        # Print cost every at intervals 10 times or as many iterations if < 10
        if i% math.ceil(num_iters/10) == 0:
            w_history.append(w)
            print(f"Iteration {i:4}: Cost {J_history[-1]:8.2f}    ")

    return w, b, J_history, w_history #return w and J,w history for graphing

```

Now let's run the gradient descent algorithm above to learn the parameters for our dataset.

```

[ ]: # Initialize fitting parameters. Recall that the shape of w is (n,)
initial_w = 0.
initial_b = 0.

# Some gradient descent settings
iterations = 1500
alpha = 0.01

w,b,_,_ = gradient_descent(x_train ,y_train, initial_w, initial_b,
                           compute_cost, compute_gradient, alpha, iterations)
print("w,b found by gradient descent:", w, b)

```

**Expected Output:**

```

w, b found by gradient descent
1.16636235 -3.63029143940436

```

## 7 Exercise 3

Now use the final parameters from gradient descent to plot the linear fit.

Recall that we can get the prediction for a single example  $f(x^{(i)}) = wx^{(i)} + b$ .

To calculate the predictions on the entire dataset, we can loop through all the training examples and calculate the prediction for each example. Write your code below.

```

[ ]: m = x_train.shape[0]
predicted = np.zeros(m)

### START YOUR CODE HERE ###

### END YOUR CODE HERE ###

```

**7.1 Now plot the predicted values to see the linear fit.**

```

[ ]: # Plot the linear fit

### START YOUR CODE HERE ###

```



```

### END YOUR CODE HERE ###

# Create a scatter plot of the data.

### START YOUR CODE HERE ###

### END YOUR CODE HERE ###

# Set the title
plt.title("Profits vs. Population per city")
# Set the y-axis label
plt.ylabel('Profit in $10,000')
# Set the x-axis label
plt.xlabel('Population of City in 10,000s')

```

## 8 Exercise 4

Your final values of  $w, b$  can also be used to make predictions on profits. Let's predict what the profit would be in areas of 35,000 and 70,000 people.

- The model takes in population of a city in 10,000s as input.
- Therefore, 35,000 people can be translated into an input to the model as `np.array([3.5])`
- Similarly, 70,000 people can be translated into an input to the model as `np.array([7.])`

```

[ ]: ### START YOUR CODE HERE ###

### END YOUR CODE HERE ###

print('For population = 35,000, we predict a profit of $%.2f' %_
      ↪(predict1*10000))

### START YOUR CODE HERE ###

### END YOUR CODE HERE ###

print('For population = 70,000, we predict a profit of $%.2f' %_
      ↪(predict2*10000))

```

### Expected Output:

For population = 35,000, we predict a profit of  
\$4519.77

For population = 70,000, we predict a profit of  
\$45342.45